



| |
|---|
| Experiment No. 6 |
| Apply Boosting Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance: |
| Date of Submission: |



Aim: Apply Boosting algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Apply Boosting algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

Theory:

Suppose that as a patient, you have certain symptoms. Instead of consulting one doctor, you choose to consult several. Suppose you assign weights to the value or worth of each doctor's diagnosis, based on the accuracies of previous diagnosis they have made. The final diagnosis is then a combination of the weighted diagnosis. This is the essence behind boosting.

Algorithm: Adaboost- A boosting algorithm—create an ensemble of classifiers. Each one gives a weighted vote.

Input:

- D , a set of d class labelled training tuples
- k , the number of rounds (one classifier is generated per round)
- a classification learning scheme

Output: A composite model

Method

1. Initialize the weight of each tuple in D is $1/d$
2. For $i=1$ to k do // for each round
3. Sample D with replacement according to the tuple weights to obtain D_i
4. Use training set D_i to derive a model M_i
5. Compute $\text{error}(M_i)$, the error rate of M_i
6. $\text{Error}(M_i) = \sum w_j \cdot \text{err}(X_j)$
7. If $\text{Error}(M_i) > 0.5$ then
8. Go back to step 3 and try again
9. endif
10. for each tuple in D_i that was correctly classified do
11. Multiply the weight of the tuple by $\text{error}(M_i)/(1-\text{error}(M_i))$
12. Normalize the weight of each tuple
13. end for

To use the ensemble to classify tuple X



1. Initialize the weight of each class to 0
2. for $i=1$ to k do // for each classifier
3. $w_i = \log((1 - \text{error}(M_i)) / \text{error}(M_i))$ // weight of the classifiers vote
4. $C = M_i(X)$ // get class prediction for X from M_i
5. Add w_i to weight for class C
6. end for
7. Return the class with the largest weight.

Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.



capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Conclusion:

1. In the conducted experiment, the XGBoost boosting algorithm yielded the following results: an accuracy of 87%, precision of 88%, recall of 94%, and an f1-score of 91%.

2. The application of the XGBoost algorithm to the Adult Income dataset has demonstrated its reliability and capacity to handle intricate data. This analysis clearly showcases XGBoost's effectiveness in boosting weaker models and optimizing predictive accuracy, particularly in the context of income level predictions, which holds significance in various socio-economic applications.

3. When comparing the outcomes of applying boosting and random forest algorithms to the Adult Census Income Dataset, it's crucial to consider the trade-offs. Boosting typically offers superior predictive accuracy, especially for complex datasets, albeit potentially sacrificing some interpretability. Conversely, random forests maintain competitive accuracy while retaining better interpretability and a higher resistance to overfitting.

```

import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import
train_test_split, cross_val_score, KFold, GridSearchCV
from sklearn.metrics import
confusion_matrix, classification_report, accuracy_score

dataset = pd.read_csv("/content/drive/MyDrive/dataset/adult.csv")

print(dataset.isnull().sum())
print(dataset.dtypes)

age          0
workclass    0
fnlwgt       0
education    0
educational-num  0
marital-status  0
occupation   0
relationship 0
race         0
gender       0
capital-gain 0
capital-loss 0
hours-per-week 0
native-country 0
income       0
dtype: int64
age          int64
workclass    object
fnlwgt       int64
education    object
educational-num  int64
marital-status  object
occupation   object
relationship  object
race         object
gender       object
capital-gain  int64
capital-loss  int64
hours-per-week  int64
native-country  object

```

```
income          object
dtype: object
```

```
dataset.head()
```

| | age | workclass | fnlwgt | education | educational-num | marital- |
|---|-----|-----------|--------|--------------|-----------------|--------------|
| 0 | 25 | Private | 226802 | 11th | 7 | Never- |
| 1 | 38 | Private | 89814 | HS-grad | 9 | Married-civ- |
| 2 | 28 | Local-gov | 336951 | Assoc-acdm | 12 | Married-civ- |
| 3 | 44 | Private | 160323 | Some-college | 10 | Married-civ- |
| 4 | 18 | ? | 103497 | Some-college | 10 | Never- |

| | occupation | relationship | race | gender | capital-gain |
|---|-------------------|--------------|-------|--------|--------------|
| 0 | Machine-op-inspct | Own-child | Black | Male | 0 |
| 1 | Farming-fishing | Husband | White | Male | 0 |
| 2 | Protective-serv | Husband | White | Male | 0 |
| 3 | Machine-op-inspct | Husband | Black | Male | 7688 |
| 4 | ? | Own-child | White | Female | 0 |

| | hours-per-week | native-country | income |
|---|----------------|----------------|--------|
| 0 | 40 | United-States | <=50K |
| 1 | 50 | United-States | <=50K |
| 2 | 40 | United-States | >50K |
| 3 | 40 | United-States | >50K |
| 4 | 30 | United-States | <=50K |

```
#removing '?' containing rows
```

```
dataset = dataset[(dataset != '?').all(axis=1)]
```

```
#label the income objects as 0 and 1
```

```
dataset['income']=dataset['income'].map({'<=50K': 0, '>50K': 1})
```

```
<ipython-input-18-39ed73805135>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation:
```


```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#
```

```

returning-a-view-versus-a-copy
dataset['income']=dataset['income'].map({'<=50K': 0, '>50K': 1})

sns.catplot(x='educational-
num',y='income',data=dataset,kind='bar',height=6)
plt.show()

```



```

#explore which country do most people belong
plt.figure(figsize=(38,14))
sns.countplot(x='native-country',data=dataset)
plt.show()

```



```
dataset['marital-status']=dataset['marital-status'].map({'Married-civ-spouse':'Married', 'Divorced':'Single', 'Never-married':'Single', 'Separated':'Single', 'Widowed':'Single', 'Married-spouse-absent':'Married', 'Married-AF-spouse':'Married'})
```

```
for column in dataset:
    enc=LabelEncoder()
    if dataset.dtypes[column]==np.object:
        dataset[column]=enc.fit_transform(dataset[column])
```

<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
if dataset.dtypes[column]==np.object:
```

<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.

Deprecated in NumPy 1.20; for more details and guidance:

<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
if dataset.dtypes[column]==np.object:
```

<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is safe.

Deprecated in NumPy 1.20; for more details and guidance:

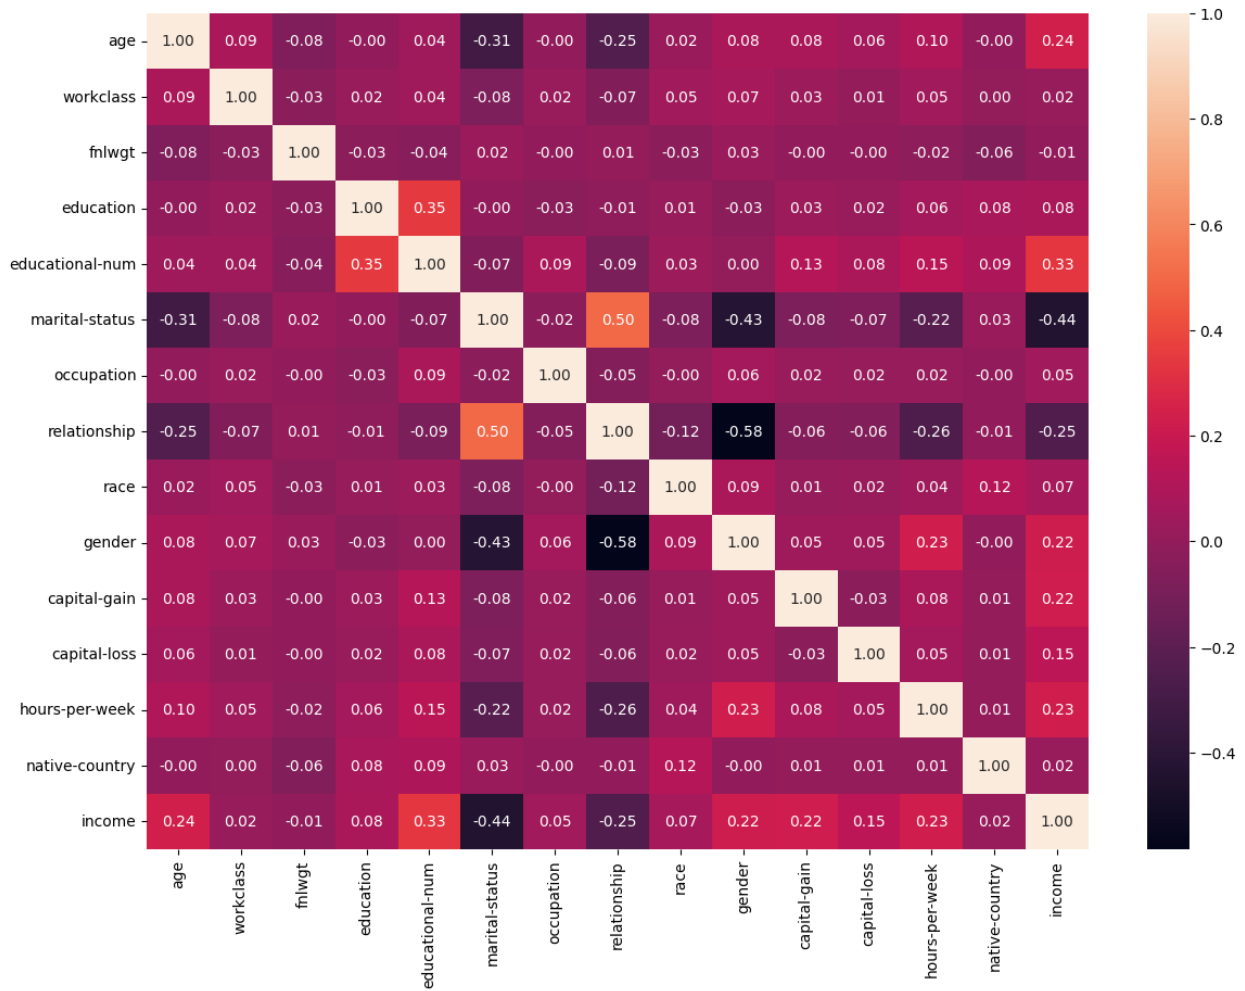
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
if dataset.dtypes[column]==np.object:
```

<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is a deprecated alias for the builtin `object`. To silence this warning, use `object` by itself. Doing this will not modify any behavior and is


```
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:
<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is
a deprecated alias for the builtin `object`. To silence this warning,
use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:
<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is
a deprecated alias for the builtin `object`. To silence this warning,
use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:
<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is
a deprecated alias for the builtin `object`. To silence this warning,
use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:
<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is
a deprecated alias for the builtin `object`. To silence this warning,
use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:
<ipython-input-24-5d7d7fe4d7c0>:3: DeprecationWarning: `np.object` is
a deprecated alias for the builtin `object`. To silence this warning,
use `object` by itself. Doing this will not modify any behavior and is
safe.
Deprecated in NumPy 1.20; for more details and guidance:
https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations
    if dataset.dtypes[column]==np.object:

plt.figure(figsize=(14,10))
sns.heatmap(dataset.corr(),annot=True,fmt='.2f')
plt.show()
```



```
dataset=dataset.drop(['relationship','education'],axis=1)
dataset=dataset.drop(['occupation','fnlwgt','native-country'],axis=1)
print(dataset.head())
```

| | age | workclass | educational-num | marital-status | race | gender | \ |
|---|-----|-----------|-----------------|----------------|------|--------|---|
| 0 | 25 | 2 | 7 | 1 | 2 | 1 | |
| 1 | 38 | 2 | 9 | 0 | 4 | 1 | |
| 2 | 28 | 1 | 12 | 0 | 4 | 1 | |
| 3 | 44 | 2 | 10 | 0 | 2 | 1 | |
| 5 | 34 | 2 | 6 | 1 | 4 | 1 | |

| | capital-gain | capital-loss | hours-per-week | income |
|---|--------------|--------------|----------------|--------|
| 0 | 0 | 0 | 40 | 0 |
| 1 | 0 | 0 | 50 | 0 |
| 2 | 0 | 0 | 40 | 1 |
| 3 | 7688 | 0 | 40 | 1 |
| 5 | 0 | 0 | 30 | 0 |

```

X=dataset.iloc[:,0:-1]
y=dataset.iloc[:,-1]
print(X.head())
print(y.head())
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.33,shuf
fle=False)

```

| | age | workclass | educational-num | marital-status | race | gender | \ |
|---|-----|-----------|-----------------|----------------|------|--------|---|
| 0 | 25 | 2 | 7 | 1 | 2 | 1 | |
| 1 | 38 | 2 | 9 | 0 | 4 | 1 | |
| 2 | 28 | 1 | 12 | 0 | 4 | 1 | |
| 3 | 44 | 2 | 10 | 0 | 2 | 1 | |
| 5 | 34 | 2 | 6 | 1 | 4 | 1 | |

| | capital-gain | capital-loss | hours-per-week |
|---|--------------|--------------|----------------|
| 0 | 0 | 0 | 40 |
| 1 | 0 | 0 | 50 |
| 2 | 0 | 0 | 40 |
| 3 | 7688 | 0 | 40 |
| 5 | 0 | 0 | 30 |

| | |
|---|---|
| 0 | 0 |
| 1 | 0 |
| 2 | 1 |
| 3 | 1 |
| 5 | 0 |

Name: income, dtype: int64

```

from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train=sc.fit_transform(X_train)
X_test=sc.transform(X_test)

```

```

from xgboost import XGBClassifier
classifier=XGBClassifier()
classifier.fit(X_train,y_train)

```

```

XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
monotone_constraints=None,

```

```
multi_strategy=None, n_estimators=None, n_jobs=None,
num_parallel_tree=None, random_state=None, ...)
```

```
from sklearn.metrics import confusion_matrix, accuracy_score ,
classification_report
y_pred=classifier.predict(X_test)
cm=confusion_matrix(y_test,y_pred)
print(cm)
accuracy_score(y_test,y_pred)
```

```
[[10532  638]
 [ 1375 2379]]
```

```
0.8651165907263468
```

```
print(classification_report(y_test , y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.88 | 0.94 | 0.91 | 11170 |
| 1 | 0.79 | 0.63 | 0.70 | 3754 |
| accuracy | | | 0.87 | 14924 |
| macro avg | 0.84 | 0.79 | 0.81 | 14924 |
| weighted avg | 0.86 | 0.87 | 0.86 | 14924 |

```
from sklearn.model_selection import cross_val_score
accuracies=cross_val_score(estimator=classifier,X=X_train,y=y_train,cv
=10)
print('Accuracy: {:.2f} Standard Deviation:
{:.2f}'.format(accuracies.mean()*100,accuracies.std()*100))
```

```
Accuracy: 85.85 Standard Deviation: 0.36
```